

Untitled1

August 8, 2024

```
[1]: import cv2
import numpy as np
from keras.datasets import cifar10
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
from keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense , Conv2D , MaxPooling2D ,Dropout, Flatten ,
↳BatchNormalization
from keras.regularizers import l2
from keras.optimizers import Adam
from keras.callbacks import ReduceLROnPlateau, EarlyStopping
from keras.models import load_model
import warnings
warnings.filterwarnings('ignore')
```

```
[2]: (X_train ,y_train),(X_test , y_test) = cifar10.load_data()
```

```
[3]: X_train ,X_valid ,y_train , y_valid = train_test_split(X_train ,y_train ,
↳test_size =0.1,random_state =0)
```

```
[4]: print('Train Image Shape : ', X_train.shape)
print('Train labels : ',y_train.shape)

print('Valid Image shape : ',X_valid.shape)
print('valid Labels : ', y_valid.shape)

print('test Image : ',X_test.shape)
print('test labels : ',y_test.shape)
```

```
Train Image Shape : (45000, 32, 32, 3)
Train labels : (45000, 1)
Valid Image shape : (5000, 32, 32, 3)
valid Labels : (5000, 1)
test Image : (10000, 32, 32, 3)
test labels : (10000, 1)
```

```
[5]: class_name = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
plt.figure(figsize=(15,15))

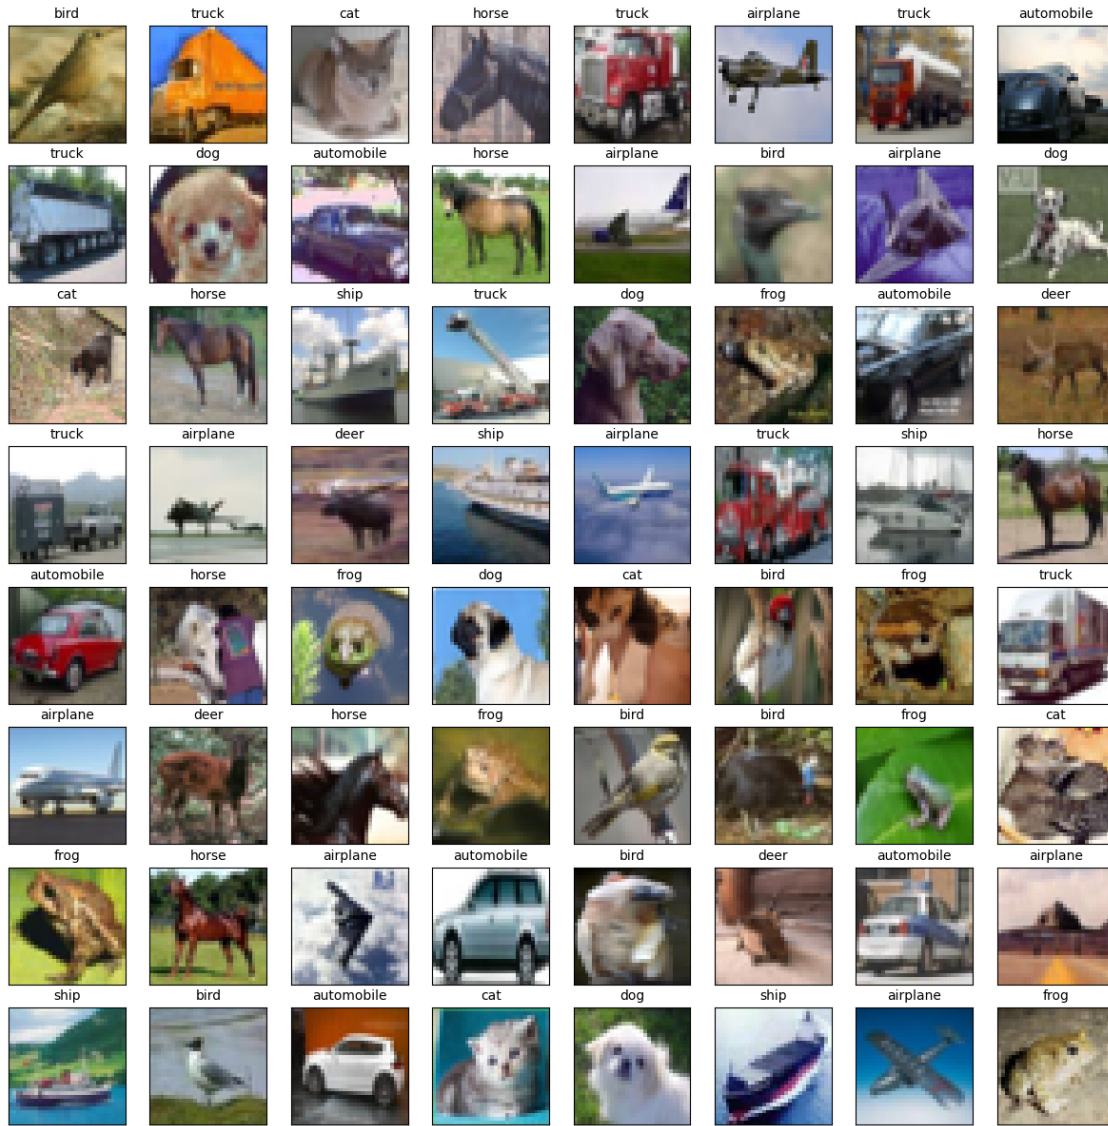
for i in range(64):

    plt.subplot(8,8,i+1)
    plt.xticks([])
    plt.yticks([])
    plt.grid(False)

    plt.imshow(X_train[i])

    plt.title(class_name[y_train[i][0]],fontsize =10)

plt.show()
```



#Normalization of Image

```
[6]: X_train = X_train.astype('float32')
      X_valid = X_valid.astype('float32')
      X_test = X_test.astype('float32')
```

```
[7]: #Calculate the mean and Standard deviation of the training image
```

```
[8]: mean = np.mean(X_train)
      std = np.std(X_train)
```

```
[9]: X_train = (X_train-mean)/(std+1e-7) # 1e-7 add for prevent the division by 0
      X_test = (X_test -mean)/(std+1e-7)
```

```
X_valid = (X_valid - mean)/(std+1e-7)
```

1 ONE - HOT ENCODING OF LABELS

`#class_names`: This is a list that maps the integer labels to human-readable names. `#Each index in this list corresponds to a class label in the one-hot encoded vector. #For instance, class_names[3] is 'cat', so the one-hot vector [0, 0, 0, 1, 0, 0, 0, 0, 0, 0] corresponds to the label 'cat'.`

```
[10]: y_train = to_categorical(y_train , 10)
      y_valid = to_categorical(y_valid , 10)
      y_test = to_categorical(y_test , 10)
```

2 Data Augmentation

is a powerful technique to enhance your dataset and improve model performance.

By applying various transformations to your data, you make your model more robust and capable of handling a wider range of real-world scenarios

Why Use Data Augmentation?

Increase Dataset Size: Helps in cases where you have a limited amount of data. By generating augmented versions of your existing data, you effectively increase the dataset size. `#####`
Improve Model Generalization: Helps the model generalize better to new, unseen data by providing more varied examples during training. `#####`
Reduce Overfitting: By introducing variations, the model is less likely to memorize the training data, which reduces overfitting.

The choice of data augmentation techniques often depends on the specific characteristics of the dataset and the problem at hand. The CIFAR-10 dataset comprises small color images of objects from 10 different classes. Given the nature of these images, some augmentation techniques are more applicable than others:

```
[11]: # Data augmentation
data_generator = ImageDataGenerator(
    # Rotate images randomly by up to 15 degrees
    rotation_range=15,

    # Shift images horizontally by up to 12% of their width
    width_shift_range=0.12,

    # Shift images vertically by up to 12% of their height
    height_shift_range=0.12,

    # Randomly flip images horizontally
    horizontal_flip=True,

    # Zoom images in by up to 10%
```

```

zoom_range=0.1,

# Change brightness by up to 10%
brightness_range=[0.9,1.1],

# Shear intensity (shear angle in counter-clockwise direction in degrees)
shear_range=10,

# Channel shift intensity
channel_shift_range=0.1,
)

```

3 Define CNN Model Architecture

The network begins with a pair of Conv2D layers, each with 32 filters of size 3x3. This is followed by a Batch Normalization layer which accelerates training and provides some level of regularization, helping to prevent overfitting.

The pairs of Conv2D layers are followed by a MaxPooling2D layer, which reduces the spatial dimensions (height and width), effectively providing a form of translation invariance and reducing computational complexity. This is followed by a Dropout layer that randomly sets a fraction (0.2 for the first dropout layer) of the input units to 0 at each update during training, helping to prevent overfitting.

This pattern of two Conv2D layers, followed by a Batch Normalization layer, a MaxPooling2D layer, and a Dropout layer, repeats three more times. The number of filters in the Conv2D layers doubles with each repetition, starting from 32 and going up to 64, 128, and then 256. This increasing pattern helps the network to learn more complex features at each level. The dropout rate also increases at each step, from 0.2 to 0.5.

After the convolutional and pooling layers, a Flatten layer is used to convert the 2D outputs of the preceding layer into a 1D vector.

Finally, a Dense (or fully connected) layer is used for classification. It has 10 units, each representing one of the 10 classes of the CIFAR-10 dataset, and a softmax activation function is used to convert the outputs to probability scores for each class.

```

[12]: model = Sequential()

# Set the weight decay value for L2 regularization
weight_decay = 0.0001

# first Convolutional layer
model.add(Conv2D(filters=32, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay), input_shape=X_train.
    ↪shape[1:]))

# add batch normalization layer
model.add(BatchNormalization())

```

```

# Second Convolutional layer (similar to first)
model.add(Conv2D(filters=32, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))
# add batch normalization layer
model.add(BatchNormalization())

# adding first max pooling layer
model.add(MaxPooling2D(pool_size=(2,2)))

# add drop out layer
model.add(Dropout(rate=0.2))

# adding the third and fourth convolutional layers with 64 filters
model.add(Conv2D(filters=64, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))
model.add(BatchNormalization())
model.add(Conv2D(filters=64, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))
model.add(BatchNormalization())

# Add the second max pooling layer and increase dropout rate to 0.3
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(rate=0.3))

# adding the fifth and sixth convolutional layers with 128 filters
model.add(Conv2D(filters=128, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))
model.add(BatchNormalization())
model.add(Conv2D(filters=128, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))
model.add(BatchNormalization())

# Add the third max pooling layer and increase dropout rate to 0.4
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(rate=0.4))

# adding the seventh and eighth convolutional layers with 256 filters

model.add(Conv2D(filters=256, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))
model.add(BatchNormalization())
model.add(Conv2D(filters=256, kernel_size=(3,3), padding='same',
    ↪activation='relu', kernel_regularizer=l2(weight_decay)))

```

```

model.add(BatchNormalization())

# Add the fourth max pooling layer and increase dropout rate to 0.5
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(rate=0.5))

# Flatten
model.add(Flatten())

model.add(Dense(10,activation='softmax'))

```

```
[13]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	896
batch_normalization (BatchNormalization)	(None, 32, 32, 32)	128
conv2d_1 (Conv2D)	(None, 32, 32, 32)	9,248
batch_normalization_1 (BatchNormalization)	(None, 32, 32, 32)	128
max_pooling2d (MaxPooling2D)	(None, 16, 16, 32)	0
dropout (Dropout)	(None, 16, 16, 32)	0
conv2d_2 (Conv2D)	(None, 16, 16, 64)	18,496
batch_normalization_2 (BatchNormalization)	(None, 16, 16, 64)	256
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36,928
batch_normalization_3 (BatchNormalization)	(None, 16, 16, 64)	256
max_pooling2d_1 (MaxPooling2D)	(None, 8, 8, 64)	0
dropout_1 (Dropout)	(None, 8, 8, 64)	0

conv2d_4 (Conv2D)	(None , 8, 8, 128)	73,856
batch_normalization_4 (BatchNormalization)	(None , 8, 8, 128)	512
conv2d_5 (Conv2D)	(None , 8, 8, 128)	147,584
batch_normalization_5 (BatchNormalization)	(None , 8, 8, 128)	512
max_pooling2d_2 (MaxPooling2D)	(None , 4, 4, 128)	0
dropout_2 (Dropout)	(None , 4, 4, 128)	0
conv2d_6 (Conv2D)	(None , 4, 4, 256)	295,168
batch_normalization_6 (BatchNormalization)	(None , 4, 4, 256)	1,024
conv2d_7 (Conv2D)	(None , 4, 4, 256)	590,080
batch_normalization_7 (BatchNormalization)	(None , 4, 4, 256)	1,024
max_pooling2d_3 (MaxPooling2D)	(None , 2, 2, 256)	0
dropout_3 (Dropout)	(None , 2, 2, 256)	0
flatten (Flatten)	(None , 1024)	0
dense (Dense)	(None , 10)	10,250

Total params: 1,186,346 (4.53 MB)

Trainable params: 1,184,426 (4.52 MB)

Non-trainable params: 1,920 (7.50 KB)

4 Training the CNN model

The `ReduceLROnPlateau` callback is used to reduce the learning rate by half (`factor=0.5`) whenever the validation loss does not improve for 10 consecutive epochs. This helps to adjust the learning rate dynamically, allowing the model to get closer to the global minimum of the loss function when progress has plateaued. This strategy

can improve the convergence of the training process.

The EarlyStopping callback is employed to monitor the validation loss and halt the training process when there hasn't been any improvement for a certain number of epochs, ensuring that the model doesn't waste computational resources and time. Furthermore, this callback restores the best weights from the training process, ensuring we conclude with the optimal model configuration from the epochs.

```
[14]: x_batch, y_batch = next(data_generator.flow(X_train, y_train,
    ↪batch_size=batch_size))
print(x_batch.shape, y_batch.shape)
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[14], line 1
----> 1 x_batch, y_batch = next(data_generator.flow(X_train, y_train,
    ↪batch_size=batch_size))
      2 print(x_batch.shape, y_batch.shape)

NameError: name 'batch_size' is not defined
```

```
[16]: # Set the batch size for the training
batch_size = 64

# Set the maximum number of epochs for the training
epochs = 300

# Define the optimizer (Adam)
optimizer = Adam()

# Compile the model with the defined optimizer, loss function, and metrics
model.compile(optimizer=optimizer, loss='categorical_crossentropy',
    ↪metrics=['accuracy'])

# Add ReduceLROnPlateau callback
reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=10,
    ↪min_lr=0.00001)

# Add EarlyStopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=40,
    ↪restore_best_weights=True, verbose=1)

# Fit the model on the training data
model.fit(data_generator.flow(X_train, y_train, batch_size=batch_size),
    epochs=epochs,
    validation_data=(X_valid, y_valid),
    callbacks=[reduce_lr, early_stopping],
```

```
verbose=1)
```

Epoch 1/300

704/704 238s 334ms/step -

accuracy: 0.3041 - loss: 2.6676 - val_accuracy: 0.4736 - val_loss: 1.7924 -

learning_rate: 0.0010

Epoch 2/300

704/704 193s 273ms/step -

accuracy: 0.4650 - loss: 1.7757 - val_accuracy: 0.5748 - val_loss: 1.3070 -

learning_rate: 0.0010

Epoch 3/300

704/704 188s 267ms/step -

accuracy: 0.5513 - loss: 1.4684 - val_accuracy: 0.6454 - val_loss: 1.1327 -

learning_rate: 0.0010

Epoch 4/300

704/704 200s 283ms/step -

accuracy: 0.6011 - loss: 1.3191 - val_accuracy: 0.6650 - val_loss: 1.0865 -

learning_rate: 0.0010

Epoch 5/300

704/704 195s 277ms/step -

accuracy: 0.6349 - loss: 1.2150 - val_accuracy: 0.7018 - val_loss: 1.0288 -

learning_rate: 0.0010

Epoch 6/300

704/704 205s 291ms/step -

accuracy: 0.6561 - loss: 1.1821 - val_accuracy: 0.7064 - val_loss: 1.0679 -

learning_rate: 0.0010

Epoch 7/300

704/704 193s 274ms/step -

accuracy: 0.6800 - loss: 1.1257 - val_accuracy: 0.6870 - val_loss: 1.2464 -

learning_rate: 0.0010

Epoch 8/300

704/704 222s 315ms/step -

accuracy: 0.6868 - loss: 1.1209 - val_accuracy: 0.7512 - val_loss: 0.9348 -

learning_rate: 0.0010

Epoch 9/300

704/704 197s 279ms/step -

accuracy: 0.7086 - loss: 1.0665 - val_accuracy: 0.7304 - val_loss: 1.0037 -

learning_rate: 0.0010

Epoch 10/300

704/704 203s 288ms/step -

accuracy: 0.7189 - loss: 1.0679 - val_accuracy: 0.7660 - val_loss: 0.9290 -

learning_rate: 0.0010

Epoch 11/300

704/704 211s 300ms/step -

accuracy: 0.7225 - loss: 1.0547 - val_accuracy: 0.7820 - val_loss: 0.8818 -

learning_rate: 0.0010

Epoch 12/300

704/704 217s 308ms/step -

accuracy: 0.7316 - loss: 1.0452 - val_accuracy: 0.7742 - val_loss: 0.9115 -
 learning_rate: 0.0010
 Epoch 13/300
 704/704 209s 297ms/step -
 accuracy: 0.7438 - loss: 1.0155 - val_accuracy: 0.7908 - val_loss: 0.8716 -
 learning_rate: 0.0010
 Epoch 14/300
 704/704 227s 322ms/step -
 accuracy: 0.7476 - loss: 1.0029 - val_accuracy: 0.7944 - val_loss: 0.8753 -
 learning_rate: 0.0010
 Epoch 15/300
 704/704 213s 301ms/step -
 accuracy: 0.7548 - loss: 0.9897 - val_accuracy: 0.7762 - val_loss: 0.9498 -
 learning_rate: 0.0010
 Epoch 16/300
 704/704 215s 305ms/step -
 accuracy: 0.7570 - loss: 0.9891 - val_accuracy: 0.7822 - val_loss: 0.8936 -
 learning_rate: 0.0010
 Epoch 17/300
 704/704 207s 293ms/step -
 accuracy: 0.7662 - loss: 0.9575 - val_accuracy: 0.7840 - val_loss: 0.9399 -
 learning_rate: 0.0010
 Epoch 18/300
 704/704 207s 294ms/step -
 accuracy: 0.7672 - loss: 0.9664 - val_accuracy: 0.8146 - val_loss: 0.8485 -
 learning_rate: 0.0010
 Epoch 19/300
 704/704 198s 281ms/step -
 accuracy: 0.7794 - loss: 0.9243 - val_accuracy: 0.8158 - val_loss: 0.8319 -
 learning_rate: 0.0010
 Epoch 20/300
 704/704 218s 310ms/step -
 accuracy: 0.7785 - loss: 0.9294 - val_accuracy: 0.8174 - val_loss: 0.8391 -
 learning_rate: 0.0010
 Epoch 21/300
 704/704 230s 326ms/step -
 accuracy: 0.7817 - loss: 0.9054 - val_accuracy: 0.7918 - val_loss: 0.8881 -
 learning_rate: 0.0010
 Epoch 22/300
 704/704 215s 305ms/step -
 accuracy: 0.7824 - loss: 0.9092 - val_accuracy: 0.7988 - val_loss: 0.9094 -
 learning_rate: 0.0010
 Epoch 23/300
 704/704 217s 307ms/step -
 accuracy: 0.7896 - loss: 0.8930 - val_accuracy: 0.7938 - val_loss: 0.8955 -
 learning_rate: 0.0010
 Epoch 24/300
 704/704 226s 320ms/step -

accuracy: 0.7945 - loss: 0.8901 - val_accuracy: 0.8134 - val_loss: 0.8520 -
 learning_rate: 0.0010
 Epoch 25/300
 704/704 220s 312ms/step -
 accuracy: 0.8014 - loss: 0.8602 - val_accuracy: 0.8288 - val_loss: 0.7839 -
 learning_rate: 0.0010
 Epoch 26/300
 704/704 218s 309ms/step -
 accuracy: 0.8027 - loss: 0.8577 - val_accuracy: 0.8294 - val_loss: 0.7845 -
 learning_rate: 0.0010
 Epoch 27/300
 704/704 219s 311ms/step -
 accuracy: 0.7994 - loss: 0.8634 - val_accuracy: 0.8242 - val_loss: 0.8207 -
 learning_rate: 0.0010
 Epoch 28/300
 704/704 217s 308ms/step -
 accuracy: 0.8021 - loss: 0.8539 - val_accuracy: 0.8352 - val_loss: 0.7633 -
 learning_rate: 0.0010
 Epoch 29/300
 704/704 224s 317ms/step -
 accuracy: 0.8058 - loss: 0.8391 - val_accuracy: 0.8174 - val_loss: 0.8196 -
 learning_rate: 0.0010
 Epoch 30/300
 704/704 217s 308ms/step -
 accuracy: 0.8020 - loss: 0.8495 - val_accuracy: 0.8266 - val_loss: 0.8085 -
 learning_rate: 0.0010
 Epoch 31/300
 704/704 217s 308ms/step -
 accuracy: 0.8104 - loss: 0.8281 - val_accuracy: 0.8320 - val_loss: 0.7780 -
 learning_rate: 0.0010
 Epoch 32/300
 704/704 219s 311ms/step -
 accuracy: 0.8065 - loss: 0.8429 - val_accuracy: 0.8310 - val_loss: 0.7722 -
 learning_rate: 0.0010
 Epoch 33/300
 704/704 220s 312ms/step -
 accuracy: 0.8117 - loss: 0.8247 - val_accuracy: 0.8318 - val_loss: 0.7791 -
 learning_rate: 0.0010
 Epoch 34/300
 704/704 212s 300ms/step -
 accuracy: 0.8132 - loss: 0.8184 - val_accuracy: 0.8474 - val_loss: 0.7269 -
 learning_rate: 0.0010
 Epoch 35/300
 704/704 230s 327ms/step -
 accuracy: 0.8126 - loss: 0.8169 - val_accuracy: 0.8140 - val_loss: 0.8086 -
 learning_rate: 0.0010
 Epoch 36/300
 704/704 234s 332ms/step -

accuracy: 0.8190 - loss: 0.8080 - val_accuracy: 0.8288 - val_loss: 0.7843 -
 learning_rate: 0.0010
 Epoch 37/300
 704/704 206s 292ms/step -
 accuracy: 0.8175 - loss: 0.8130 - val_accuracy: 0.8368 - val_loss: 0.7696 -
 learning_rate: 0.0010
 Epoch 38/300
 704/704 207s 294ms/step -
 accuracy: 0.8204 - loss: 0.7988 - val_accuracy: 0.8450 - val_loss: 0.7274 -
 learning_rate: 0.0010
 Epoch 39/300
 704/704 216s 307ms/step -
 accuracy: 0.8167 - loss: 0.8060 - val_accuracy: 0.8530 - val_loss: 0.7068 -
 learning_rate: 0.0010
 Epoch 40/300
 704/704 192s 272ms/step -
 accuracy: 0.8189 - loss: 0.7895 - val_accuracy: 0.8398 - val_loss: 0.7440 -
 learning_rate: 0.0010
 Epoch 41/300
 704/704 217s 308ms/step -
 accuracy: 0.8209 - loss: 0.7926 - val_accuracy: 0.8546 - val_loss: 0.6921 -
 learning_rate: 0.0010
 Epoch 42/300
 704/704 239s 339ms/step -
 accuracy: 0.8196 - loss: 0.7865 - val_accuracy: 0.8380 - val_loss: 0.7318 -
 learning_rate: 0.0010
 Epoch 43/300
 704/704 220s 312ms/step -
 accuracy: 0.8224 - loss: 0.7809 - val_accuracy: 0.8570 - val_loss: 0.6928 -
 learning_rate: 0.0010
 Epoch 44/300
 704/704 213s 302ms/step -
 accuracy: 0.8220 - loss: 0.7845 - val_accuracy: 0.8350 - val_loss: 0.7819 -
 learning_rate: 0.0010
 Epoch 45/300
 704/704 197s 280ms/step -
 accuracy: 0.8252 - loss: 0.7789 - val_accuracy: 0.8538 - val_loss: 0.6973 -
 learning_rate: 0.0010
 Epoch 46/300
 704/704 196s 278ms/step -
 accuracy: 0.8233 - loss: 0.7845 - val_accuracy: 0.8460 - val_loss: 0.7322 -
 learning_rate: 0.0010
 Epoch 47/300
 704/704 205s 291ms/step -
 accuracy: 0.8263 - loss: 0.7728 - val_accuracy: 0.8552 - val_loss: 0.7019 -
 learning_rate: 0.0010
 Epoch 48/300
 704/704 207s 294ms/step -

accuracy: 0.8264 - loss: 0.7742 - val_accuracy: 0.8380 - val_loss: 0.7416 -
 learning_rate: 0.0010
 Epoch 49/300
 704/704 211s 299ms/step -
 accuracy: 0.8290 - loss: 0.7570 - val_accuracy: 0.8512 - val_loss: 0.7007 -
 learning_rate: 0.0010
 Epoch 50/300
 704/704 209s 297ms/step -
 accuracy: 0.8266 - loss: 0.7740 - val_accuracy: 0.8390 - val_loss: 0.7608 -
 learning_rate: 0.0010
 Epoch 51/300
 704/704 230s 326ms/step -
 accuracy: 0.8295 - loss: 0.7664 - val_accuracy: 0.8456 - val_loss: 0.7248 -
 learning_rate: 0.0010
 Epoch 52/300
 704/704 265s 376ms/step -
 accuracy: 0.8400 - loss: 0.7318 - val_accuracy: 0.8678 - val_loss: 0.6342 -
 learning_rate: 5.0000e-04
 Epoch 53/300
 704/704 264s 374ms/step -
 accuracy: 0.8514 - loss: 0.6886 - val_accuracy: 0.8700 - val_loss: 0.6298 -
 learning_rate: 5.0000e-04
 Epoch 54/300
 704/704 209s 296ms/step -
 accuracy: 0.8529 - loss: 0.6705 - val_accuracy: 0.8704 - val_loss: 0.6146 -
 learning_rate: 5.0000e-04
 Epoch 55/300
 704/704 206s 292ms/step -
 accuracy: 0.8557 - loss: 0.6546 - val_accuracy: 0.8664 - val_loss: 0.6237 -
 learning_rate: 5.0000e-04
 Epoch 56/300
 704/704 199s 283ms/step -
 accuracy: 0.8554 - loss: 0.6581 - val_accuracy: 0.8720 - val_loss: 0.6067 -
 learning_rate: 5.0000e-04
 Epoch 57/300
 704/704 196s 279ms/step -
 accuracy: 0.8578 - loss: 0.6373 - val_accuracy: 0.8746 - val_loss: 0.5855 -
 learning_rate: 5.0000e-04
 Epoch 58/300
 704/704 199s 282ms/step -
 accuracy: 0.8572 - loss: 0.6366 - val_accuracy: 0.8688 - val_loss: 0.6040 -
 learning_rate: 5.0000e-04
 Epoch 59/300
 704/704 311s 442ms/step -
 accuracy: 0.8586 - loss: 0.6304 - val_accuracy: 0.8596 - val_loss: 0.6374 -
 learning_rate: 5.0000e-04
 Epoch 60/300
 704/704 401s 568ms/step -

accuracy: 0.8571 - loss: 0.6237 - val_accuracy: 0.8688 - val_loss: 0.6159 -
 learning_rate: 5.0000e-04
 Epoch 61/300
 704/704 296s 420ms/step -
 accuracy: 0.8598 - loss: 0.6233 - val_accuracy: 0.8690 - val_loss: 0.5969 -
 learning_rate: 5.0000e-04
 Epoch 62/300
 704/704 240s 341ms/step -
 accuracy: 0.8620 - loss: 0.6116 - val_accuracy: 0.8732 - val_loss: 0.5929 -
 learning_rate: 5.0000e-04
 Epoch 63/300
 704/704 255s 362ms/step -
 accuracy: 0.8632 - loss: 0.6026 - val_accuracy: 0.8782 - val_loss: 0.5756 -
 learning_rate: 5.0000e-04
 Epoch 64/300
 704/704 356s 505ms/step -
 accuracy: 0.8645 - loss: 0.6010 - val_accuracy: 0.8844 - val_loss: 0.5431 -
 learning_rate: 5.0000e-04
 Epoch 65/300
 704/704 266s 377ms/step -
 accuracy: 0.8648 - loss: 0.5991 - val_accuracy: 0.8756 - val_loss: 0.5817 -
 learning_rate: 5.0000e-04
 Epoch 66/300
 704/704 245s 348ms/step -
 accuracy: 0.8652 - loss: 0.5991 - val_accuracy: 0.8742 - val_loss: 0.5846 -
 learning_rate: 5.0000e-04
 Epoch 67/300
 704/704 231s 328ms/step -
 accuracy: 0.8640 - loss: 0.5915 - val_accuracy: 0.8690 - val_loss: 0.5822 -
 learning_rate: 5.0000e-04
 Epoch 68/300
 704/704 246s 349ms/step -
 accuracy: 0.8657 - loss: 0.5881 - val_accuracy: 0.8774 - val_loss: 0.5693 -
 learning_rate: 5.0000e-04
 Epoch 69/300
 704/704 236s 335ms/step -
 accuracy: 0.8673 - loss: 0.5918 - val_accuracy: 0.8738 - val_loss: 0.5787 -
 learning_rate: 5.0000e-04
 Epoch 70/300
 704/704 239s 338ms/step -
 accuracy: 0.8666 - loss: 0.5850 - val_accuracy: 0.8726 - val_loss: 0.5922 -
 learning_rate: 5.0000e-04
 Epoch 71/300
 704/704 232s 329ms/step -
 accuracy: 0.8642 - loss: 0.5855 - val_accuracy: 0.8728 - val_loss: 0.5840 -
 learning_rate: 5.0000e-04
 Epoch 72/300
 704/704 241s 342ms/step -

accuracy: 0.8639 - loss: 0.5868 - val_accuracy: 0.8814 - val_loss: 0.5457 -
 learning_rate: 5.0000e-04
 Epoch 73/300
 704/704 236s 335ms/step -
 accuracy: 0.8652 - loss: 0.5867 - val_accuracy: 0.8664 - val_loss: 0.5951 -
 learning_rate: 5.0000e-04
 Epoch 74/300
 704/704 235s 333ms/step -
 accuracy: 0.8660 - loss: 0.5853 - val_accuracy: 0.8716 - val_loss: 0.5889 -
 learning_rate: 5.0000e-04
 Epoch 75/300
 704/704 255s 362ms/step -
 accuracy: 0.8712 - loss: 0.5638 - val_accuracy: 0.8852 - val_loss: 0.5415 -
 learning_rate: 2.5000e-04
 Epoch 76/300
 704/704 255s 361ms/step -
 accuracy: 0.8795 - loss: 0.5450 - val_accuracy: 0.8826 - val_loss: 0.5328 -
 learning_rate: 2.5000e-04
 Epoch 77/300
 704/704 273s 388ms/step -
 accuracy: 0.8810 - loss: 0.5315 - val_accuracy: 0.8910 - val_loss: 0.5198 -
 learning_rate: 2.5000e-04
 Epoch 78/300
 704/704 290s 411ms/step -
 accuracy: 0.8799 - loss: 0.5271 - val_accuracy: 0.8892 - val_loss: 0.5250 -
 learning_rate: 2.5000e-04
 Epoch 79/300
 704/704 262s 372ms/step -
 accuracy: 0.8804 - loss: 0.5216 - val_accuracy: 0.8824 - val_loss: 0.5407 -
 learning_rate: 2.5000e-04
 Epoch 80/300
 704/704 246s 349ms/step -
 accuracy: 0.8832 - loss: 0.5142 - val_accuracy: 0.8832 - val_loss: 0.5255 -
 learning_rate: 2.5000e-04
 Epoch 81/300
 704/704 266s 377ms/step -
 accuracy: 0.8842 - loss: 0.5090 - val_accuracy: 0.8892 - val_loss: 0.5079 -
 learning_rate: 2.5000e-04
 Epoch 82/300
 704/704 241s 341ms/step -
 accuracy: 0.8838 - loss: 0.5095 - val_accuracy: 0.8840 - val_loss: 0.5228 -
 learning_rate: 2.5000e-04
 Epoch 83/300
 704/704 241s 340ms/step -
 accuracy: 0.8870 - loss: 0.4990 - val_accuracy: 0.8872 - val_loss: 0.5118 -
 learning_rate: 2.5000e-04
 Epoch 84/300
 704/704 239s 340ms/step -

accuracy: 0.8886 - loss: 0.4911 - val_accuracy: 0.8812 - val_loss: 0.5378 -
 learning_rate: 2.5000e-04
 Epoch 85/300
 704/704 242s 343ms/step -
 accuracy: 0.8863 - loss: 0.4859 - val_accuracy: 0.8928 - val_loss: 0.5024 -
 learning_rate: 2.5000e-04
 Epoch 86/300
 704/704 267s 380ms/step -
 accuracy: 0.8866 - loss: 0.4916 - val_accuracy: 0.8920 - val_loss: 0.5026 -
 learning_rate: 2.5000e-04
 Epoch 87/300
 704/704 250s 355ms/step -
 accuracy: 0.8865 - loss: 0.4946 - val_accuracy: 0.8946 - val_loss: 0.4759 -
 learning_rate: 2.5000e-04
 Epoch 88/300
 704/704 229s 325ms/step -
 accuracy: 0.8825 - loss: 0.4983 - val_accuracy: 0.8924 - val_loss: 0.4971 -
 learning_rate: 2.5000e-04
 Epoch 89/300
 704/704 230s 326ms/step -
 accuracy: 0.8889 - loss: 0.4829 - val_accuracy: 0.8840 - val_loss: 0.5184 -
 learning_rate: 2.5000e-04
 Epoch 90/300
 704/704 250s 354ms/step -
 accuracy: 0.8910 - loss: 0.4789 - val_accuracy: 0.8952 - val_loss: 0.4753 -
 learning_rate: 2.5000e-04
 Epoch 91/300
 704/704 282s 400ms/step -
 accuracy: 0.8890 - loss: 0.4829 - val_accuracy: 0.8912 - val_loss: 0.4944 -
 learning_rate: 2.5000e-04
 Epoch 92/300
 704/704 274s 389ms/step -
 accuracy: 0.8895 - loss: 0.4807 - val_accuracy: 0.8906 - val_loss: 0.5134 -
 learning_rate: 2.5000e-04
 Epoch 93/300
 704/704 229s 326ms/step -
 accuracy: 0.8905 - loss: 0.4742 - val_accuracy: 0.8974 - val_loss: 0.4687 -
 learning_rate: 2.5000e-04
 Epoch 94/300
 704/704 265s 376ms/step -
 accuracy: 0.8913 - loss: 0.4728 - val_accuracy: 0.8944 - val_loss: 0.4761 -
 learning_rate: 2.5000e-04
 Epoch 95/300
 704/704 261s 371ms/step -
 accuracy: 0.8908 - loss: 0.4767 - val_accuracy: 0.8874 - val_loss: 0.4912 -
 learning_rate: 2.5000e-04
 Epoch 96/300
 704/704 225s 320ms/step -

accuracy: 0.8907 - loss: 0.4703 - val_accuracy: 0.8918 - val_loss: 0.4967 -
 learning_rate: 2.5000e-04
 Epoch 97/300
 704/704 242s 344ms/step -
 accuracy: 0.8894 - loss: 0.4733 - val_accuracy: 0.8910 - val_loss: 0.4878 -
 learning_rate: 2.5000e-04
 Epoch 98/300
 704/704 242s 343ms/step -
 accuracy: 0.8895 - loss: 0.4720 - val_accuracy: 0.8898 - val_loss: 0.5073 -
 learning_rate: 2.5000e-04
 Epoch 99/300
 704/704 243s 344ms/step -
 accuracy: 0.8896 - loss: 0.4733 - val_accuracy: 0.8896 - val_loss: 0.4911 -
 learning_rate: 2.5000e-04
 Epoch 100/300
 704/704 205s 291ms/step -
 accuracy: 0.8885 - loss: 0.4731 - val_accuracy: 0.8858 - val_loss: 0.5016 -
 learning_rate: 2.5000e-04
 Epoch 101/300
 704/704 221s 314ms/step -
 accuracy: 0.8927 - loss: 0.4663 - val_accuracy: 0.8946 - val_loss: 0.4724 -
 learning_rate: 2.5000e-04
 Epoch 102/300
 704/704 206s 292ms/step -
 accuracy: 0.8908 - loss: 0.4641 - val_accuracy: 0.8926 - val_loss: 0.4803 -
 learning_rate: 2.5000e-04
 Epoch 103/300
 704/704 206s 292ms/step -
 accuracy: 0.8897 - loss: 0.4663 - val_accuracy: 0.8848 - val_loss: 0.5063 -
 learning_rate: 2.5000e-04
 Epoch 104/300
 704/704 203s 288ms/step -
 accuracy: 0.8932 - loss: 0.4596 - val_accuracy: 0.8956 - val_loss: 0.4675 -
 learning_rate: 1.2500e-04
 Epoch 105/300
 704/704 256s 364ms/step -
 accuracy: 0.8985 - loss: 0.4428 - val_accuracy: 0.8984 - val_loss: 0.4804 -
 learning_rate: 1.2500e-04
 Epoch 106/300
 704/704 267s 380ms/step -
 accuracy: 0.8987 - loss: 0.4375 - val_accuracy: 0.8950 - val_loss: 0.4905 -
 learning_rate: 1.2500e-04
 Epoch 107/300
 704/704 240s 341ms/step -
 accuracy: 0.8992 - loss: 0.4323 - val_accuracy: 0.8986 - val_loss: 0.4660 -
 learning_rate: 1.2500e-04
 Epoch 108/300
 704/704 193s 273ms/step -

accuracy: 0.8988 - loss: 0.4337 - val_accuracy: 0.8974 - val_loss: 0.4628 -
 learning_rate: 1.2500e-04
 Epoch 109/300
 704/704 201s 285ms/step -
 accuracy: 0.9011 - loss: 0.4258 - val_accuracy: 0.8970 - val_loss: 0.4682 -
 learning_rate: 1.2500e-04
 Epoch 110/300
 704/704 206s 293ms/step -
 accuracy: 0.8989 - loss: 0.4264 - val_accuracy: 0.8950 - val_loss: 0.4755 -
 learning_rate: 1.2500e-04
 Epoch 111/300
 704/704 206s 293ms/step -
 accuracy: 0.9018 - loss: 0.4191 - val_accuracy: 0.8992 - val_loss: 0.4570 -
 learning_rate: 1.2500e-04
 Epoch 112/300
 704/704 207s 294ms/step -
 accuracy: 0.9039 - loss: 0.4157 - val_accuracy: 0.9024 - val_loss: 0.4463 -
 learning_rate: 1.2500e-04
 Epoch 113/300
 704/704 194s 275ms/step -
 accuracy: 0.9023 - loss: 0.4144 - val_accuracy: 0.9020 - val_loss: 0.4422 -
 learning_rate: 1.2500e-04
 Epoch 114/300
 704/704 243s 346ms/step -
 accuracy: 0.9032 - loss: 0.4171 - val_accuracy: 0.9016 - val_loss: 0.4446 -
 learning_rate: 1.2500e-04
 Epoch 115/300
 704/704 220s 312ms/step -
 accuracy: 0.9057 - loss: 0.4072 - val_accuracy: 0.8952 - val_loss: 0.4677 -
 learning_rate: 1.2500e-04
 Epoch 116/300
 704/704 238s 338ms/step -
 accuracy: 0.9053 - loss: 0.4125 - val_accuracy: 0.9046 - val_loss: 0.4332 -
 learning_rate: 1.2500e-04
 Epoch 117/300
 704/704 226s 321ms/step -
 accuracy: 0.9042 - loss: 0.4076 - val_accuracy: 0.8992 - val_loss: 0.4549 -
 learning_rate: 1.2500e-04
 Epoch 118/300
 704/704 264s 375ms/step -
 accuracy: 0.9068 - loss: 0.4018 - val_accuracy: 0.8998 - val_loss: 0.4573 -
 learning_rate: 1.2500e-04
 Epoch 119/300
 704/704 249s 354ms/step -
 accuracy: 0.9052 - loss: 0.4081 - val_accuracy: 0.9034 - val_loss: 0.4361 -
 learning_rate: 1.2500e-04
 Epoch 120/300
 704/704 268s 380ms/step -

accuracy: 0.9071 - loss: 0.3989 - val_accuracy: 0.8982 - val_loss: 0.4467 -
 learning_rate: 1.2500e-04
 Epoch 121/300
 704/704 351s 498ms/step -
 accuracy: 0.9047 - loss: 0.4042 - val_accuracy: 0.9022 - val_loss: 0.4304 -
 learning_rate: 1.2500e-04
 Epoch 122/300
 704/704 324s 460ms/step -
 accuracy: 0.9110 - loss: 0.3926 - val_accuracy: 0.9044 - val_loss: 0.4397 -
 learning_rate: 1.2500e-04
 Epoch 123/300
 704/704 269s 382ms/step -
 accuracy: 0.9049 - loss: 0.4036 - val_accuracy: 0.9022 - val_loss: 0.4355 -
 learning_rate: 1.2500e-04
 Epoch 124/300
 704/704 264s 375ms/step -
 accuracy: 0.9076 - loss: 0.3996 - val_accuracy: 0.9012 - val_loss: 0.4366 -
 learning_rate: 1.2500e-04
 Epoch 125/300
 704/704 208s 295ms/step -
 accuracy: 0.9035 - loss: 0.4015 - val_accuracy: 0.8984 - val_loss: 0.4441 -
 learning_rate: 1.2500e-04
 Epoch 126/300
 704/704 209s 296ms/step -
 accuracy: 0.9103 - loss: 0.3894 - val_accuracy: 0.8972 - val_loss: 0.4479 -
 learning_rate: 1.2500e-04
 Epoch 127/300
 704/704 194s 275ms/step -
 accuracy: 0.9065 - loss: 0.3960 - val_accuracy: 0.9008 - val_loss: 0.4334 -
 learning_rate: 1.2500e-04
 Epoch 128/300
 704/704 193s 274ms/step -
 accuracy: 0.9099 - loss: 0.3884 - val_accuracy: 0.8982 - val_loss: 0.4401 -
 learning_rate: 1.2500e-04
 Epoch 129/300
 704/704 203s 287ms/step -
 accuracy: 0.9083 - loss: 0.3914 - val_accuracy: 0.9056 - val_loss: 0.4210 -
 learning_rate: 1.2500e-04
 Epoch 130/300
 704/704 200s 284ms/step -
 accuracy: 0.9076 - loss: 0.3898 - val_accuracy: 0.9050 - val_loss: 0.4351 -
 learning_rate: 1.2500e-04
 Epoch 131/300
 704/704 204s 289ms/step -
 accuracy: 0.9077 - loss: 0.3907 - val_accuracy: 0.8974 - val_loss: 0.4500 -
 learning_rate: 1.2500e-04
 Epoch 132/300
 704/704 202s 287ms/step -

accuracy: 0.9099 - loss: 0.3932 - val_accuracy: 0.8996 - val_loss: 0.4451 -
 learning_rate: 1.2500e-04
 Epoch 133/300
 704/704 205s 291ms/step -
 accuracy: 0.9047 - loss: 0.3993 - val_accuracy: 0.9010 - val_loss: 0.4428 -
 learning_rate: 1.2500e-04
 Epoch 134/300
 704/704 206s 292ms/step -
 accuracy: 0.9088 - loss: 0.3879 - val_accuracy: 0.9002 - val_loss: 0.4460 -
 learning_rate: 1.2500e-04
 Epoch 135/300
 704/704 266s 378ms/step -
 accuracy: 0.9049 - loss: 0.3935 - val_accuracy: 0.9038 - val_loss: 0.4337 -
 learning_rate: 1.2500e-04
 Epoch 136/300
 704/704 266s 377ms/step -
 accuracy: 0.9091 - loss: 0.3840 - val_accuracy: 0.9014 - val_loss: 0.4336 -
 learning_rate: 1.2500e-04
 Epoch 137/300
 704/704 226s 321ms/step -
 accuracy: 0.9098 - loss: 0.3834 - val_accuracy: 0.9004 - val_loss: 0.4427 -
 learning_rate: 1.2500e-04
 Epoch 138/300
 704/704 195s 277ms/step -
 accuracy: 0.9097 - loss: 0.3807 - val_accuracy: 0.8992 - val_loss: 0.4445 -
 learning_rate: 1.2500e-04
 Epoch 139/300
 704/704 192s 273ms/step -
 accuracy: 0.9073 - loss: 0.3891 - val_accuracy: 0.8974 - val_loss: 0.4433 -
 learning_rate: 1.2500e-04
 Epoch 140/300
 704/704 193s 274ms/step -
 accuracy: 0.9073 - loss: 0.3870 - val_accuracy: 0.8986 - val_loss: 0.4444 -
 learning_rate: 6.2500e-05
 Epoch 141/300
 704/704 192s 272ms/step -
 accuracy: 0.9132 - loss: 0.3696 - val_accuracy: 0.9060 - val_loss: 0.4116 -
 learning_rate: 6.2500e-05
 Epoch 142/300
 704/704 203s 288ms/step -
 accuracy: 0.9101 - loss: 0.3702 - val_accuracy: 0.9014 - val_loss: 0.4293 -
 learning_rate: 6.2500e-05
 Epoch 143/300
 704/704 197s 279ms/step -
 accuracy: 0.9151 - loss: 0.3671 - val_accuracy: 0.9056 - val_loss: 0.4208 -
 learning_rate: 6.2500e-05
 Epoch 144/300
 704/704 199s 283ms/step -

accuracy: 0.9176 - loss: 0.3617 - val_accuracy: 0.9040 - val_loss: 0.4257 -
 learning_rate: 6.2500e-05
 Epoch 145/300
 704/704 200s 284ms/step -
 accuracy: 0.9156 - loss: 0.3649 - val_accuracy: 0.9066 - val_loss: 0.4215 -
 learning_rate: 6.2500e-05
 Epoch 146/300
 704/704 197s 280ms/step -
 accuracy: 0.9158 - loss: 0.3572 - val_accuracy: 0.9060 - val_loss: 0.4192 -
 learning_rate: 6.2500e-05
 Epoch 147/300
 704/704 234s 332ms/step -
 accuracy: 0.9143 - loss: 0.3613 - val_accuracy: 0.9050 - val_loss: 0.4194 -
 learning_rate: 6.2500e-05
 Epoch 148/300
 704/704 243s 344ms/step -
 accuracy: 0.9149 - loss: 0.3640 - val_accuracy: 0.9054 - val_loss: 0.4104 -
 learning_rate: 6.2500e-05
 Epoch 149/300
 704/704 214s 303ms/step -
 accuracy: 0.9162 - loss: 0.3580 - val_accuracy: 0.9070 - val_loss: 0.4113 -
 learning_rate: 6.2500e-05
 Epoch 150/300
 704/704 203s 289ms/step -
 accuracy: 0.9189 - loss: 0.3567 - val_accuracy: 0.9074 - val_loss: 0.4143 -
 learning_rate: 6.2500e-05
 Epoch 151/300
 704/704 236s 334ms/step -
 accuracy: 0.9151 - loss: 0.3601 - val_accuracy: 0.9078 - val_loss: 0.4104 -
 learning_rate: 6.2500e-05
 Epoch 152/300
 704/704 248s 352ms/step -
 accuracy: 0.9163 - loss: 0.3544 - val_accuracy: 0.9058 - val_loss: 0.4170 -
 learning_rate: 6.2500e-05
 Epoch 153/300
 704/704 207s 294ms/step -
 accuracy: 0.9187 - loss: 0.3527 - val_accuracy: 0.9064 - val_loss: 0.4223 -
 learning_rate: 6.2500e-05
 Epoch 154/300
 704/704 221s 313ms/step -
 accuracy: 0.9165 - loss: 0.3542 - val_accuracy: 0.9066 - val_loss: 0.4141 -
 learning_rate: 6.2500e-05
 Epoch 155/300
 704/704 220s 312ms/step -
 accuracy: 0.9158 - loss: 0.3539 - val_accuracy: 0.9088 - val_loss: 0.4016 -
 learning_rate: 6.2500e-05
 Epoch 156/300
 704/704 211s 300ms/step -

accuracy: 0.9151 - loss: 0.3554 - val_accuracy: 0.9058 - val_loss: 0.4120 -
 learning_rate: 6.2500e-05
 Epoch 157/300
 704/704 206s 293ms/step -
 accuracy: 0.9159 - loss: 0.3558 - val_accuracy: 0.9060 - val_loss: 0.4113 -
 learning_rate: 6.2500e-05
 Epoch 158/300
 704/704 267s 379ms/step -
 accuracy: 0.9181 - loss: 0.3525 - val_accuracy: 0.9044 - val_loss: 0.4078 -
 learning_rate: 6.2500e-05
 Epoch 159/300
 704/704 256s 362ms/step -
 accuracy: 0.9202 - loss: 0.3431 - val_accuracy: 0.9068 - val_loss: 0.4116 -
 learning_rate: 6.2500e-05
 Epoch 160/300
 704/704 260s 369ms/step -
 accuracy: 0.9164 - loss: 0.3521 - val_accuracy: 0.9044 - val_loss: 0.4153 -
 learning_rate: 6.2500e-05
 Epoch 161/300
 704/704 241s 342ms/step -
 accuracy: 0.9169 - loss: 0.3491 - val_accuracy: 0.9084 - val_loss: 0.4121 -
 learning_rate: 6.2500e-05
 Epoch 162/300
 704/704 219s 311ms/step -
 accuracy: 0.9203 - loss: 0.3442 - val_accuracy: 0.9074 - val_loss: 0.4094 -
 learning_rate: 6.2500e-05
 Epoch 163/300
 704/704 214s 303ms/step -
 accuracy: 0.9158 - loss: 0.3501 - val_accuracy: 0.9062 - val_loss: 0.4046 -
 learning_rate: 6.2500e-05
 Epoch 164/300
 704/704 223s 317ms/step -
 accuracy: 0.9202 - loss: 0.3447 - val_accuracy: 0.9062 - val_loss: 0.4033 -
 learning_rate: 6.2500e-05
 Epoch 165/300
 704/704 228s 324ms/step -
 accuracy: 0.9190 - loss: 0.3429 - val_accuracy: 0.9066 - val_loss: 0.4133 -
 learning_rate: 6.2500e-05
 Epoch 166/300
 704/704 246s 349ms/step -
 accuracy: 0.9206 - loss: 0.3394 - val_accuracy: 0.9074 - val_loss: 0.4137 -
 learning_rate: 3.1250e-05
 Epoch 167/300
 704/704 220s 312ms/step -
 accuracy: 0.9217 - loss: 0.3350 - val_accuracy: 0.9104 - val_loss: 0.4051 -
 learning_rate: 3.1250e-05
 Epoch 168/300
 704/704 221s 314ms/step -

accuracy: 0.9227 - loss: 0.3393 - val_accuracy: 0.9094 - val_loss: 0.4034 -
 learning_rate: 3.1250e-05
 Epoch 169/300
 704/704 202s 286ms/step -
 accuracy: 0.9182 - loss: 0.3417 - val_accuracy: 0.9082 - val_loss: 0.4040 -
 learning_rate: 3.1250e-05
 Epoch 170/300
 704/704 240s 341ms/step -
 accuracy: 0.9229 - loss: 0.3345 - val_accuracy: 0.9076 - val_loss: 0.4127 -
 learning_rate: 3.1250e-05
 Epoch 171/300
 704/704 259s 368ms/step -
 accuracy: 0.9190 - loss: 0.3447 - val_accuracy: 0.9070 - val_loss: 0.4044 -
 learning_rate: 3.1250e-05
 Epoch 172/300
 704/704 236s 335ms/step -
 accuracy: 0.9211 - loss: 0.3388 - val_accuracy: 0.9096 - val_loss: 0.4086 -
 learning_rate: 3.1250e-05
 Epoch 173/300
 704/704 227s 322ms/step -
 accuracy: 0.9205 - loss: 0.3370 - val_accuracy: 0.9100 - val_loss: 0.3987 -
 learning_rate: 3.1250e-05
 Epoch 174/300
 704/704 239s 339ms/step -
 accuracy: 0.9211 - loss: 0.3360 - val_accuracy: 0.9088 - val_loss: 0.4021 -
 learning_rate: 3.1250e-05
 Epoch 175/300
 704/704 228s 324ms/step -
 accuracy: 0.9227 - loss: 0.3321 - val_accuracy: 0.9108 - val_loss: 0.3932 -
 learning_rate: 3.1250e-05
 Epoch 176/300
 704/704 230s 327ms/step -
 accuracy: 0.9199 - loss: 0.3383 - val_accuracy: 0.9110 - val_loss: 0.4026 -
 learning_rate: 3.1250e-05
 Epoch 177/300
 704/704 225s 319ms/step -
 accuracy: 0.9241 - loss: 0.3296 - val_accuracy: 0.9114 - val_loss: 0.4010 -
 learning_rate: 3.1250e-05
 Epoch 178/300
 704/704 236s 335ms/step -
 accuracy: 0.9225 - loss: 0.3330 - val_accuracy: 0.9098 - val_loss: 0.4015 -
 learning_rate: 3.1250e-05
 Epoch 179/300
 704/704 240s 341ms/step -
 accuracy: 0.9261 - loss: 0.3219 - val_accuracy: 0.9090 - val_loss: 0.3958 -
 learning_rate: 3.1250e-05
 Epoch 180/300
 704/704 247s 351ms/step -

accuracy: 0.9257 - loss: 0.3271 - val_accuracy: 0.9118 - val_loss: 0.3961 -
 learning_rate: 3.1250e-05
 Epoch 181/300
 704/704 222s 315ms/step -
 accuracy: 0.9272 - loss: 0.3222 - val_accuracy: 0.9076 - val_loss: 0.4029 -
 learning_rate: 3.1250e-05
 Epoch 182/300
 704/704 223s 317ms/step -
 accuracy: 0.9231 - loss: 0.3309 - val_accuracy: 0.9082 - val_loss: 0.3987 -
 learning_rate: 3.1250e-05
 Epoch 183/300
 704/704 232s 329ms/step -
 accuracy: 0.9234 - loss: 0.3308 - val_accuracy: 0.9058 - val_loss: 0.4108 -
 learning_rate: 3.1250e-05
 Epoch 184/300
 704/704 248s 352ms/step -
 accuracy: 0.9250 - loss: 0.3268 - val_accuracy: 0.9064 - val_loss: 0.4004 -
 learning_rate: 3.1250e-05
 Epoch 185/300
 704/704 213s 303ms/step -
 accuracy: 0.9225 - loss: 0.3291 - val_accuracy: 0.9086 - val_loss: 0.3936 -
 learning_rate: 3.1250e-05
 Epoch 186/300
 704/704 202s 286ms/step -
 accuracy: 0.9231 - loss: 0.3242 - val_accuracy: 0.9084 - val_loss: 0.3984 -
 learning_rate: 1.5625e-05
 Epoch 187/300
 704/704 205s 291ms/step -
 accuracy: 0.9225 - loss: 0.3259 - val_accuracy: 0.9106 - val_loss: 0.3947 -
 learning_rate: 1.5625e-05
 Epoch 188/300
 704/704 208s 295ms/step -
 accuracy: 0.9220 - loss: 0.3303 - val_accuracy: 0.9104 - val_loss: 0.3982 -
 learning_rate: 1.5625e-05
 Epoch 189/300
 704/704 203s 288ms/step -
 accuracy: 0.9224 - loss: 0.3280 - val_accuracy: 0.9094 - val_loss: 0.3951 -
 learning_rate: 1.5625e-05
 Epoch 190/300
 704/704 209s 297ms/step -
 accuracy: 0.9265 - loss: 0.3248 - val_accuracy: 0.9104 - val_loss: 0.3977 -
 learning_rate: 1.5625e-05
 Epoch 191/300
 704/704 204s 290ms/step -
 accuracy: 0.9241 - loss: 0.3265 - val_accuracy: 0.9102 - val_loss: 0.3955 -
 learning_rate: 1.5625e-05
 Epoch 192/300
 704/704 202s 287ms/step -

accuracy: 0.9276 - loss: 0.3139 - val_accuracy: 0.9100 - val_loss: 0.3979 -
 learning_rate: 1.5625e-05
 Epoch 193/300
 704/704 206s 292ms/step -
 accuracy: 0.9226 - loss: 0.3298 - val_accuracy: 0.9110 - val_loss: 0.3955 -
 learning_rate: 1.5625e-05
 Epoch 194/300
 704/704 269s 382ms/step -
 accuracy: 0.9232 - loss: 0.3303 - val_accuracy: 0.9112 - val_loss: 0.3914 -
 learning_rate: 1.5625e-05
 Epoch 195/300
 704/704 279s 396ms/step -
 accuracy: 0.9234 - loss: 0.3233 - val_accuracy: 0.9104 - val_loss: 0.3982 -
 learning_rate: 1.5625e-05
 Epoch 196/300
 704/704 236s 334ms/step -
 accuracy: 0.9234 - loss: 0.3260 - val_accuracy: 0.9118 - val_loss: 0.3945 -
 learning_rate: 1.5625e-05
 Epoch 197/300
 704/704 201s 285ms/step -
 accuracy: 0.9251 - loss: 0.3211 - val_accuracy: 0.9118 - val_loss: 0.3924 -
 learning_rate: 1.5625e-05
 Epoch 198/300
 704/704 198s 281ms/step -
 accuracy: 0.9218 - loss: 0.3246 - val_accuracy: 0.9116 - val_loss: 0.3937 -
 learning_rate: 1.5625e-05
 Epoch 199/300
 704/704 199s 282ms/step -
 accuracy: 0.9237 - loss: 0.3251 - val_accuracy: 0.9112 - val_loss: 0.3938 -
 learning_rate: 1.5625e-05
 Epoch 200/300
 704/704 205s 290ms/step -
 accuracy: 0.9262 - loss: 0.3155 - val_accuracy: 0.9108 - val_loss: 0.3952 -
 learning_rate: 1.5625e-05
 Epoch 201/300
 704/704 201s 285ms/step -
 accuracy: 0.9260 - loss: 0.3171 - val_accuracy: 0.9100 - val_loss: 0.3971 -
 learning_rate: 1.5625e-05
 Epoch 202/300
 704/704 199s 283ms/step -
 accuracy: 0.9227 - loss: 0.3276 - val_accuracy: 0.9118 - val_loss: 0.3886 -
 learning_rate: 1.5625e-05
 Epoch 203/300
 704/704 200s 285ms/step -
 accuracy: 0.9262 - loss: 0.3141 - val_accuracy: 0.9104 - val_loss: 0.3963 -
 learning_rate: 1.5625e-05
 Epoch 204/300
 704/704 201s 285ms/step -

accuracy: 0.9275 - loss: 0.3142 - val_accuracy: 0.9090 - val_loss: 0.3980 -
 learning_rate: 1.5625e-05
 Epoch 205/300
 704/704 202s 287ms/step -
 accuracy: 0.9262 - loss: 0.3205 - val_accuracy: 0.9110 - val_loss: 0.3921 -
 learning_rate: 1.5625e-05
 Epoch 206/300
 704/704 199s 283ms/step -
 accuracy: 0.9267 - loss: 0.3159 - val_accuracy: 0.9114 - val_loss: 0.4010 -
 learning_rate: 1.5625e-05
 Epoch 207/300
 704/704 202s 286ms/step -
 accuracy: 0.9268 - loss: 0.3173 - val_accuracy: 0.9108 - val_loss: 0.3982 -
 learning_rate: 1.5625e-05
 Epoch 208/300
 704/704 199s 283ms/step -
 accuracy: 0.9272 - loss: 0.3201 - val_accuracy: 0.9120 - val_loss: 0.3949 -
 learning_rate: 1.5625e-05
 Epoch 209/300
 704/704 208s 295ms/step -
 accuracy: 0.9284 - loss: 0.3126 - val_accuracy: 0.9098 - val_loss: 0.3962 -
 learning_rate: 1.5625e-05
 Epoch 210/300
 704/704 202s 286ms/step -
 accuracy: 0.9271 - loss: 0.3129 - val_accuracy: 0.9094 - val_loss: 0.3990 -
 learning_rate: 1.5625e-05
 Epoch 211/300
 704/704 224s 317ms/step -
 accuracy: 0.9282 - loss: 0.3121 - val_accuracy: 0.9108 - val_loss: 0.3932 -
 learning_rate: 1.5625e-05
 Epoch 212/300
 704/704 281s 399ms/step -
 accuracy: 0.9252 - loss: 0.3203 - val_accuracy: 0.9102 - val_loss: 0.3941 -
 learning_rate: 1.5625e-05
 Epoch 213/300
 704/704 233s 331ms/step -
 accuracy: 0.9273 - loss: 0.3162 - val_accuracy: 0.9106 - val_loss: 0.3925 -
 learning_rate: 1.0000e-05
 Epoch 214/300
 704/704 231s 327ms/step -
 accuracy: 0.9266 - loss: 0.3153 - val_accuracy: 0.9098 - val_loss: 0.3953 -
 learning_rate: 1.0000e-05
 Epoch 215/300
 704/704 227s 322ms/step -
 accuracy: 0.9274 - loss: 0.3137 - val_accuracy: 0.9116 - val_loss: 0.3897 -
 learning_rate: 1.0000e-05
 Epoch 216/300
 704/704 228s 324ms/step -

accuracy: 0.9280 - loss: 0.3129 - val_accuracy: 0.9092 - val_loss: 0.3986 -
 learning_rate: 1.0000e-05
 Epoch 217/300
 704/704 246s 349ms/step -
 accuracy: 0.9266 - loss: 0.3163 - val_accuracy: 0.9114 - val_loss: 0.3922 -
 learning_rate: 1.0000e-05
 Epoch 218/300
 704/704 232s 329ms/step -
 accuracy: 0.9237 - loss: 0.3202 - val_accuracy: 0.9118 - val_loss: 0.3907 -
 learning_rate: 1.0000e-05
 Epoch 219/300
 704/704 230s 326ms/step -
 accuracy: 0.9278 - loss: 0.3103 - val_accuracy: 0.9098 - val_loss: 0.3958 -
 learning_rate: 1.0000e-05
 Epoch 220/300
 704/704 228s 324ms/step -
 accuracy: 0.9274 - loss: 0.3117 - val_accuracy: 0.9110 - val_loss: 0.3947 -
 learning_rate: 1.0000e-05
 Epoch 221/300
 704/704 233s 331ms/step -
 accuracy: 0.9286 - loss: 0.3104 - val_accuracy: 0.9122 - val_loss: 0.3985 -
 learning_rate: 1.0000e-05
 Epoch 222/300
 704/704 232s 330ms/step -
 accuracy: 0.9262 - loss: 0.3128 - val_accuracy: 0.9112 - val_loss: 0.3935 -
 learning_rate: 1.0000e-05
 Epoch 223/300
 704/704 235s 334ms/step -
 accuracy: 0.9276 - loss: 0.3123 - val_accuracy: 0.9102 - val_loss: 0.3945 -
 learning_rate: 1.0000e-05
 Epoch 224/300
 704/704 233s 331ms/step -
 accuracy: 0.9288 - loss: 0.3064 - val_accuracy: 0.9120 - val_loss: 0.3925 -
 learning_rate: 1.0000e-05
 Epoch 225/300
 704/704 252s 357ms/step -
 accuracy: 0.9270 - loss: 0.3150 - val_accuracy: 0.9118 - val_loss: 0.3907 -
 learning_rate: 1.0000e-05
 Epoch 226/300
 704/704 245s 348ms/step -
 accuracy: 0.9264 - loss: 0.3136 - val_accuracy: 0.9108 - val_loss: 0.3938 -
 learning_rate: 1.0000e-05
 Epoch 227/300
 704/704 243s 345ms/step -
 accuracy: 0.9264 - loss: 0.3158 - val_accuracy: 0.9104 - val_loss: 0.3918 -
 learning_rate: 1.0000e-05
 Epoch 228/300
 704/704 245s 348ms/step -

accuracy: 0.9262 - loss: 0.3146 - val_accuracy: 0.9114 - val_loss: 0.3925 -
 learning_rate: 1.0000e-05
 Epoch 229/300
 704/704 239s 339ms/step -
 accuracy: 0.9272 - loss: 0.3116 - val_accuracy: 0.9118 - val_loss: 0.3942 -
 learning_rate: 1.0000e-05
 Epoch 230/300
 704/704 247s 350ms/step -
 accuracy: 0.9253 - loss: 0.3184 - val_accuracy: 0.9120 - val_loss: 0.3934 -
 learning_rate: 1.0000e-05
 Epoch 231/300
 704/704 256s 363ms/step -
 accuracy: 0.9297 - loss: 0.3076 - val_accuracy: 0.9114 - val_loss: 0.3987 -
 learning_rate: 1.0000e-05
 Epoch 232/300
 704/704 255s 362ms/step -
 accuracy: 0.9280 - loss: 0.3078 - val_accuracy: 0.9132 - val_loss: 0.3921 -
 learning_rate: 1.0000e-05
 Epoch 233/300
 704/704 206s 293ms/step -
 accuracy: 0.9259 - loss: 0.3173 - val_accuracy: 0.9114 - val_loss: 0.3958 -
 learning_rate: 1.0000e-05
 Epoch 234/300
 704/704 205s 291ms/step -
 accuracy: 0.9288 - loss: 0.3056 - val_accuracy: 0.9128 - val_loss: 0.3904 -
 learning_rate: 1.0000e-05
 Epoch 235/300
 704/704 210s 298ms/step -
 accuracy: 0.9276 - loss: 0.3131 - val_accuracy: 0.9098 - val_loss: 0.3934 -
 learning_rate: 1.0000e-05
 Epoch 236/300
 704/704 207s 294ms/step -
 accuracy: 0.9244 - loss: 0.3194 - val_accuracy: 0.9110 - val_loss: 0.3934 -
 learning_rate: 1.0000e-05
 Epoch 237/300
 704/704 204s 290ms/step -
 accuracy: 0.9262 - loss: 0.3121 - val_accuracy: 0.9114 - val_loss: 0.3935 -
 learning_rate: 1.0000e-05
 Epoch 238/300
 704/704 204s 290ms/step -
 accuracy: 0.9272 - loss: 0.3135 - val_accuracy: 0.9104 - val_loss: 0.3940 -
 learning_rate: 1.0000e-05
 Epoch 239/300
 704/704 205s 291ms/step -
 accuracy: 0.9282 - loss: 0.3112 - val_accuracy: 0.9102 - val_loss: 0.3915 -
 learning_rate: 1.0000e-05
 Epoch 240/300
 704/704 214s 303ms/step -

```
accuracy: 0.9270 - loss: 0.3120 - val_accuracy: 0.9106 - val_loss: 0.3915 -  
learning_rate: 1.0000e-05  
Epoch 241/300  
704/704          213s 303ms/step -  
accuracy: 0.9284 - loss: 0.3111 - val_accuracy: 0.9092 - val_loss: 0.3943 -  
learning_rate: 1.0000e-05  
Epoch 242/300  
704/704          206s 292ms/step -  
accuracy: 0.9260 - loss: 0.3137 - val_accuracy: 0.9098 - val_loss: 0.3945 -  
learning_rate: 1.0000e-05  
Epoch 242: early stopping  
Restoring model weights from the end of the best epoch: 202.
```

```
[16]: <keras.src.callbacks.history.History at 0x3173f1a10>
```

```
[17]: test_loss, test_acc = model.evaluate(X_test, y_test, verbose=1)  
  
print('\nTest Accuracy:', test_acc)  
print('Test Loss:      ', test_loss)
```

```
313/313          12s 38ms/step -  
accuracy: 0.9101 - loss: 0.3975
```

```
Test Accuracy: 0.9083999991416931  
Test Loss:      0.40374499559402466
```

```
[ ]:
```